

Distributed tracing with Jaeger

Closing dev-ops gap in performance analysis

Outlines


- What is distributed tracing
- Why do we need it
- Opentracing: terms and concepts
- Jaeger implementation
- Live coding demo: write simple apps with distributed tracing libraries

Distributed tracing

- Tracing activity which crosses the barrier between processes
- End-to-end tracing technology, especially for microservices
- Collect details about events from different services
- Put them in the same context (a trace) of request execution flow
- Form the causality relationship between services
- To provide a complete picture of your distributed applications

Why distributed tracing ?

- 1 request spreads through multiple services
- A need for step-by-step instrumentation over the network
- Context passing and merging
- No-root, language-agnostic implementation across services
- Inspect your services like a browser's requests
- Service dependencies visualization

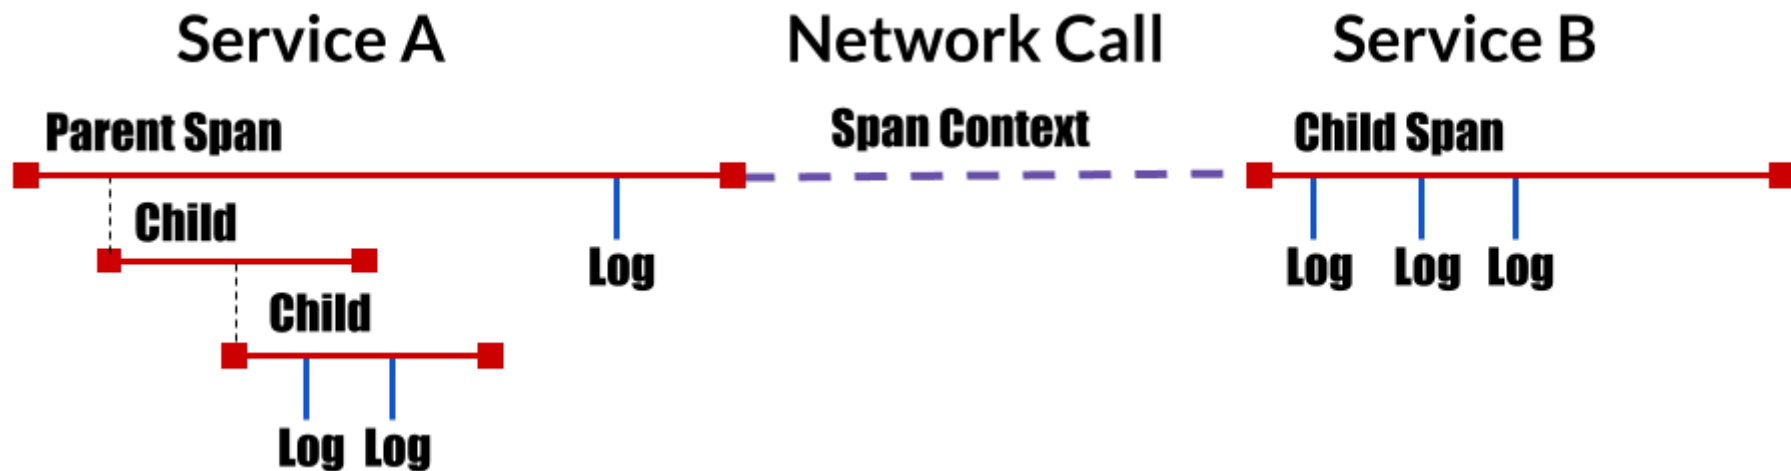


Opentracing

An standardization effort for distributed tracing system, which define:

- tracing instrumentation API:
 - Tracer initialization
 - Inject / Extract
- Wire protocol: format of data sent along side with app's data (trace context)
 - IDs
 - References
 - Tag, Log , Baggage
- Data protocol: propagation msg format , transport protocol, encodes

Opentracing terms



- A library that enables observability for another library is called an instrumentation library
- API / SDK
- Collector + agent for data gathering
- Backend tracing: for data processing and visualization

~~Three pillars~~ Triangle of observability

- Traces
- Metrics
- Logs

API vs SDK

- If you're developing a library or some other component that is intended to be consumed by a runnable binary, then you would only take a dependency on the API.
- If your artifact is a standalone process or service, then you would take a dependency on the API and the SDK.

Technology Stack

- Go backend
- Pluggable storage
 - Cassandra, Elasticsearch, memory, ...
- React/Javascript frontend
- OpenTracing Instrumentation libraries
- Integration with Kafka, Apache Flink



elasticsearch



OPENTRACING

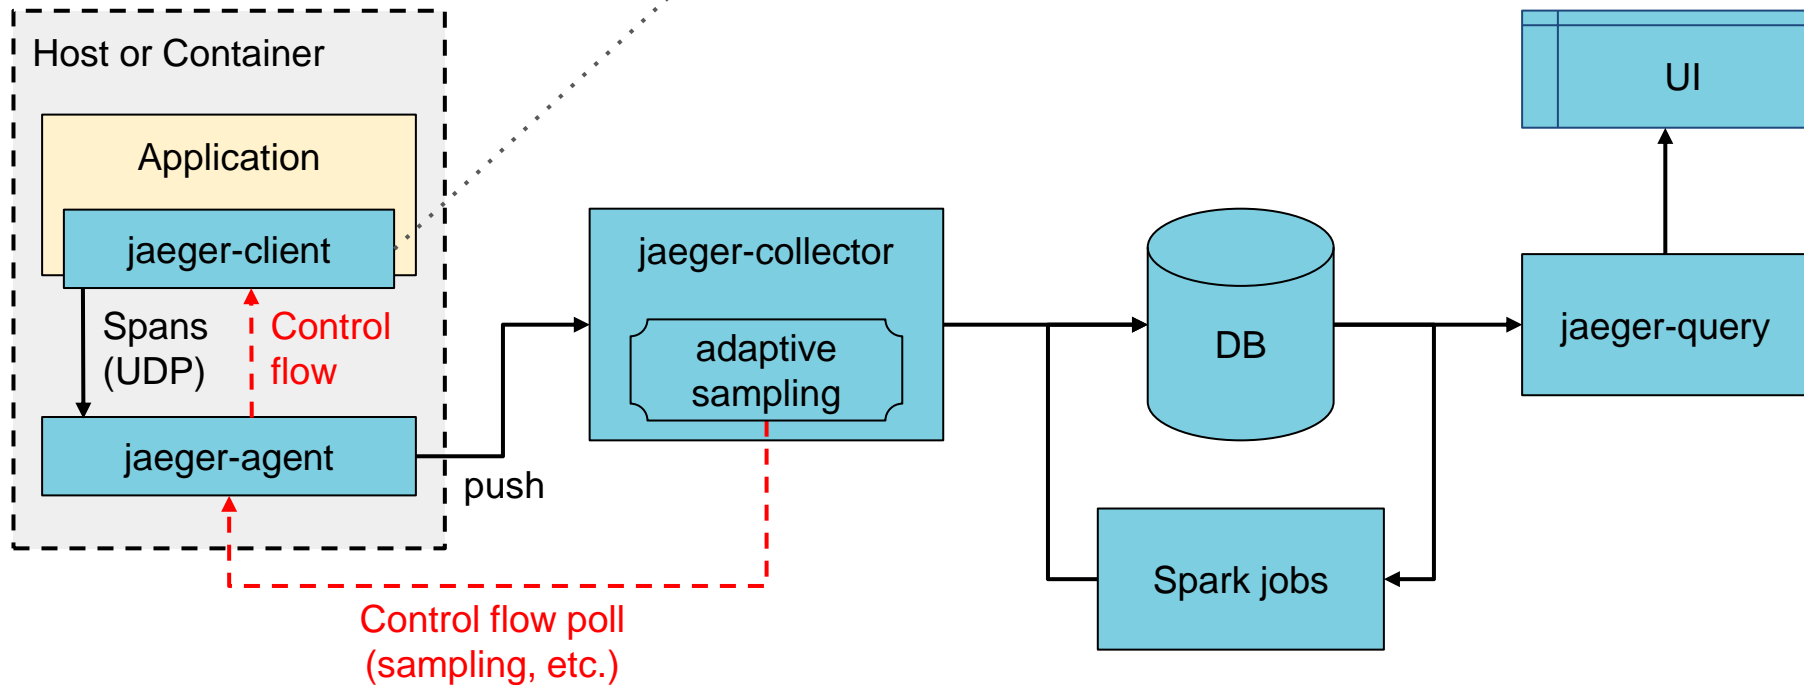


POWERED



CLOUD NATIVE
COMPUTING FOUNDATION

Apache Cassandra® is a trademark of the [Apache Software Foundation](https://www.apache.org/) in the United States and/or other countries.



Ecosystem and alternatives

Distributed backend tracing solution:

- Origin: [dapper](#)
- Opensource: [zipkin](#), [skywalking](#)
- Enterprise solutions: [lightstep](#), [stackdriver](#), [aws x-ray](#)


Distributed tracing API / SDK/ specification:

- Opentracing
- Opentelemetry
- Opencensus

Vendors

- <https://opentelemetry.io/vendors/>

Live demo



Write simple apps armed with distributed tracing

- Prepare trace initialization config
- Start a tracer
- Start a span
- Passing span inside process manually
- Use active span to automate context passing
- Tag, log, baggage
- Context propagation between services

Challenges

- Need to update codes
- Convincing developers is hard
- Performance of o11y instrumentation libraries and infrastructure
- Mature but fast changing community

Culture shift required

- Strong dev/ops collaboration
- Observability and performance mindset

Future plans

- High throughput with kafka
- Smart adaptive sampling mechanism
- Integration with existing metrics (gr, prom) to provide “context-rich” paging msg

References

- <https://opentracing.io/docs/>
- <https://www.jaegertracing.io/>
- <https://www.shkuro.com/books/2019-mastering-distributed-tracing/>
- <https://kubernetes.github.io/ingress-nginx/user-guide/third-party-addons/opentracing/>
- <https://github.com/opentracing/opentracing-python>
- <https://opentelemetry.io/docs/concepts/>
- <https://medium.com/jaegertracing/jaeger-and-opentelemetry-1846f701d9f2>